

Final Project Progress

Emma Chen, Xavier Routh, William Yeh

Please remember to prepare **3 to 5 pages of slides** for these meetings.

This will help you better introduce the paper you are working on, propose the new methodology and describe the progress achieved by your team so far.

November 13-20:

10% of the grade

Each team will sign up for a 15 minute progress meeting with the TA. Before the meeting, the students should write the progress report, 1-page PDF and email it to the TA. By this stage you should have a *partial working code for a subset of the proposed functionality* that has been tested on test cases you wrote. The functionality can be small, but it should be solid and non-trivial (original to your project). Describe what you have accomplished, including any relevant preliminary results for programs that work. Note that 30% of the overall project grade is reserved for your progress accomplished during the first month including your project proposal.

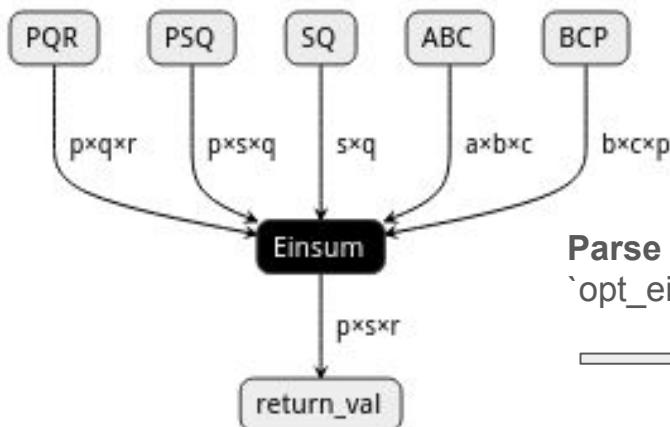
Einsum Introduction

- Einsum expressions lowered into a tree of binary operations between input tensors
- Einsum contraction paths (binary trees) in terms of FLOPs and Memory have good solutions already
- Large einsum expressions can benefit from global optimization to optimize for data layout.
- Given a contraction tree (that is already optimized for FLOPs and Memory):
- **Einsum Tree IR** is used to optimize for data layout, and to target fast primitives (transpose, GEMM, packed GEMM)

Progress 1

ONNX Expression:

pqr,psq,sq,abc,bcp->psr



Parse / Find contraction path:
`'opt_einsum.contract_path("pqr, ...->")'`

Einsum Tree IR

Constructed Tree:

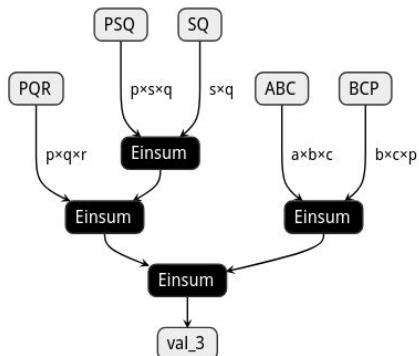
- node_3: spr,p->psr
- node_2: sqp,pqr->spr
 - node_1: sq,psq->sqp
 - sq: sq
 - psq: psq
 - pqr: pqr
- node_0: bcp,abc->p
 - bcp: bcp
 - abc: abc

Progress 2

Unoptimized Einsum Tree IR

Constructed Tree:

- node_3: spr,p->psr
 - node_2: sqp,pqr->spr
 - node_1: sq,psq->sqp
 - sq: sq
 - psq: psq
 - pqr: pqr
 - node_0: bcp,abc->p
 - bcp: bcp
 - abc: abc

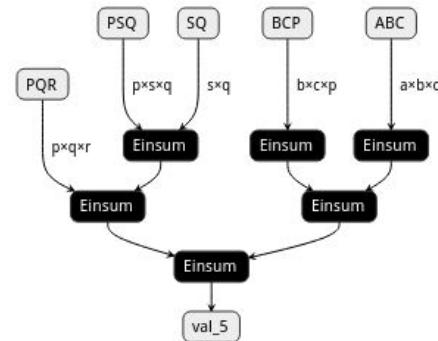


Optimization from Einsum Tree Paper



Optimized Einsum Tree IR

- node_3: psr,p->psr
 - node_2: pqr,psq->psr
 - pqr: pqr->pqr
 - pqr: pqr
 - node_1: psq,sq->psq
 - psq: psq->psq
 - psq: psq
 - sq: sq->sq
 - sq: sq
 - node_0: cbp,acb->p
 - cbp: bcp->cbp
 - bcp: bcp
 - acb: abc->acb
 - abc: abc



Results

Setup:

Personal computer CPU
(ONNX RT provider)
20 runs

Numpy Reference:
np.einsum(expr)

ONNX non-optimized:
Single onnx einsum node

ONNX optimized:
Original contraction tree

ONNX optimized2:
Einsum Tree paper optimized
contraction tree

```
def sizes():
    return {
        "a": 20,
        "b": 30,
        "c": 10,
        "p": 200,
        "q": 60,
        "r": 3,
        "s": 2,
    }
```

```
Benchmarking provider=CPUExecutionProvider, warmup=3, runs=10
NumPy reference: avg 1949.635 ms, 0.5 it/s
ONNX non-optimized: avg 1.375 ms, 727.2 it/s
ONNX optimized: avg 0.332 ms, 3013.9 it/s
ONNX optimized2: avg 0.151 ms, 6620.7 it/s
```

Next Steps

- Benchmark on GPU
- Codegen (lowering contraction path to loops + 3 primitives):
 - **Problem:** Attempts to build into ONNX fail, as onnx doesn't support packed GEMM. No data layout definitions in ONNX.
 - **Current Solution:** Leave to ONNX RT.
 - **Future Solution:** Build codegen algorithm into XLA frontend, as it rewrites einsum expressions (stableHLO). See if we can enforce intermediate data layouts this way.
- Evaluate impact of original contraction path